Part H — Computer vision & OCR

ml-ladder writeup · part H

2026-09-18

Contents

H1 — Detection, and the metric as a construction **2**
The math . . . 2

H2 — Segmentation, and averaging over the wrong pixels **5**
The math . . . 5

H3 — CTC, and what "held out" actually held out **7**
The math . . . 7

H4 — Two ways to draw the image-to-string interface **10**
The math . . . 10

H5 — The pipeline, and where its error actually lives **13**
The math . . . 13

H6 — Deleting the interface **15**
The math . . . 16

What part H settled **18**

Status: complete (H1–H6) · **Constraint:** synthetic data with exact ground truth, rendered offline; small budgets
Numbers: seeds 0–2, reference seed 0 · raw values in `data/h-metrics.json`
Reproduce: `ladder run H1` … `H6` (needs `uv sync --extra torch --extra ocr`), or `uv run python tools/writeup_figures.py H`
Hardware: H1–H5 swept on a Colab T4; H6 is arithmetic and runs anywhere

Every part before this one was handed its metric. Accuracy needs no argument — the prediction equals the label or it does not — and even part D's probes were scored by counting correct answers. Detection has no such luxury. A model emits a variable number of boxes, each with a score and a class, and *nothing* says which prediction corresponds to which object. The metric has to be built, and building it involves choices that change the answer.

So part H is about **interfaces**: between a prediction and the thing it is supposed to be a prediction *of*, and between one stage of a pipeline and the next. H1 constructs a metric from nothing. H2 discovers that averaging over the wrong pixels hides the effect it is looking for. H3 and H4 find that the interface between an image and a string can be drawn two ways with opposite failure modes. H5 measures a real two-stage pipeline and finds **99.98% of its error living in the interface between the stages**. H6 prices what it costs to delete that interface entirely, which is what the field actually did.

The data is rendered rather than downloaded, in `src/mlladder/synth.py`. Attributing a pipeline's error to a stage needs ground truth that is exactly right — which box, which character, which pixel — and real OCR corpora come with annotation noise that is itself a research topic. The fonts come from matplotlib's bundled TTFs rather than the system's, because otherwise every number here would be a property of the machine that produced it. That reproducibility is verified, not assumed: the rendered pixels are byte-identical across matplotlib 3.10/3.11 and Pillow 11.3/12.3, checked by hashing the output on two machines and now guarded by a test.

Notation: boxes are (x_1, y_1, x_2, y_2) in pixels, IoU is intersection over union, CER and WER are edit distances normalized by reference characters and words, and T is a sequence length in image columns or tokens.

H1 — Detection, and the metric as a construction

Question. Detection's score depends on three choices that have nothing to do with the detector: what overlap counts as a match, how duplicates are removed, and how a precision-recall curve becomes one number. How much do those choices move the answer?

The math

IoU is the whole foundation, and it is one ratio:

$$\mathrm{IoU}(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|}.$$

The intersection of two axis-aligned boxes is a box — the coordinate-wise max of the near corners and min of the far ones — which is why this is cheap enough to evaluate on every prediction-truth pair.

Non-max suppression is three lines and one assumption: that two boxes overlapping by more than τ_{nms} describe the same object. Sort by score, keep the best, discard everything it overlaps, repeat. The assumption is false whenever two real objects overlap that much, which is why the threshold is a dial rather than a constant.

Average precision walks the predictions in descending score order. A prediction is a true positive if it clears τ_{iou} against a truth in the same image *that nothing better has already claimed* — the claiming rule is what makes a duplicate cost something, and therefore what makes NMS worth running at all. Each prefix gives a (r, p) point, and AP is the area under the monotone envelope:

$$\text{AP} = \sum_k (r_k - r_{k-1}) \max_{\tilde{r} \geq r_k} p(\tilde{r}).$$

Taking the running maximum of precision from the right is the "interpolated" part. It matters more than it looks: a duplicate that lands *after* full recall has been reached costs nothing at all under the envelope, so NMS only pays where duplicates outrank genuine detections elsewhere.

mAP averages AP over classes — over the classes that are *present*. Scoring an absent class 0 rather than skipping it is a silent way to make a detector look bad on a dataset that never showed it that class.

What I built. A YOLO-style detector with one prediction per cell of an 8×8 grid on 64×64 scenes: objectness, a box, and a class distribution per cell. The box is parameterised relative to its cell — centre as a sigmoid offset inside it, size as a log-ratio of the image — because predicting absolute coordinates asks every cell to learn where it is, while offsets let one shared convolution work everywhere. Ground-truth boxes are assigned to the cell holding their centre, which caps recall when two centres share a cell; that happens to 0.17% of objects here and is reported rather than left implicit.

Results (3000 train / 600 test scenes, 12 epochs; medians over seeds 0–2):

NMS threshold	boxes/image	mAP@0.5	mAP@0.75
none	4.30	0.847	0.652
0.1	1.91	0.923	0.634
0.3	1.99	0.953	0.649
0.5	2.20	**0.955**	0.650
0.7	3.38	0.908	**0.672**
0.9	4.30	0.847	0.652

Ground truth averages 2.0 objects per image, which NMS at 0.3 reproduces almost exactly (1.99). `best/nms_threshold` is 0.5 on every seed.

What surprised me. The two IoU thresholds disagree about what the NMS dial is *for*. Across the sweep, mAP@0.5 moves by **0.109** and mAP@0.75 by only **0.032** — and at 0.75 the best threshold is 0.7, not 0.5. At the strict threshold, duplicate removal stops mattering because localization has become the binding constraint, and a knob that looks important at one threshold is nearly inert at the other.

The localization penalty itself is the largest number in the rung: **0.955 at IoU 0.5 against 0.650 at 0.75**, a drop of 0.316 on *unchanged predictions*. I chased that rather than accept it, and the chase is most of what I learned here. It is not a grid bug — the predicted centre is unbiased to 0.1 px,

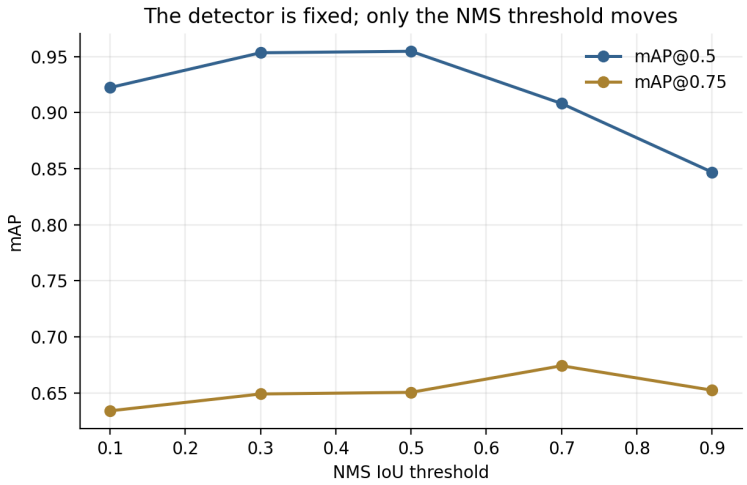


Figure 1: NMS sweep

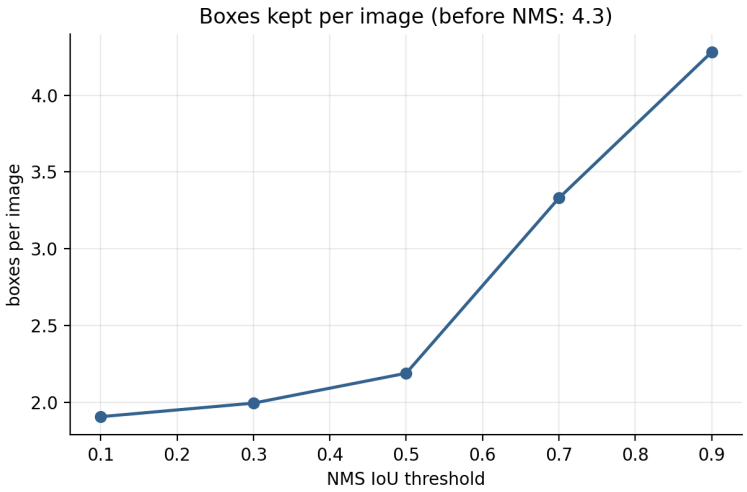


Figure 2: Boxes per image

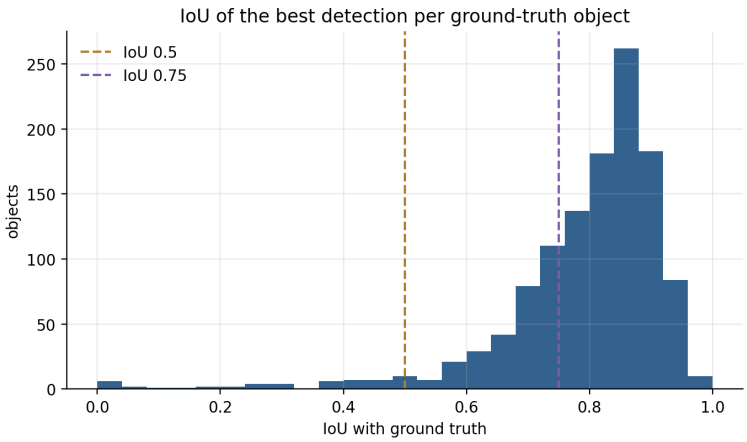


Figure 3: Matched IoU

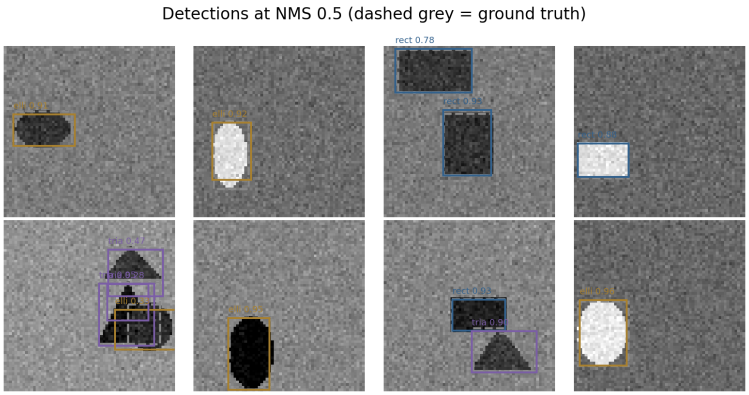


Figure 4: Detections

and feeding the encoder's own targets through the decoder round-trips to IoU > 0.99, which is now a test. It is not shape fill either: my hypothesis was that a triangle, filling half its box, would localize worst, but the width bias is identical across rectangle, ellipse and triangle (-1.46, -1.44, -1.49 px). And it is not Jensen's inequality from the log-size parameterisation, which accounts for only ~0.6 px of it. The residual is biased in log space itself (-0.067), where the *variance* (std 0.11, an 11% relative size error) dominates the bias anyway. That variance is what caps matched IoU at 0.792 mean, with 97% of objects above 0.5 but only 11% above 0.9.

Per class at the strict threshold the ordering is ellipse 0.698 $>$ rectangle 0.629 $>$ triangle 0.579, which my fill hypothesis gets exactly half right: the triangle is worst, but a rectangle *is* its box and still loses to an ellipse. I do not have an explanation I believe, and one sentence of honest ignorance beats a plausible story.

Open threads. Anchors with aspect priors, which is the standard answer to a size regression this noisy. Focal loss for the 30:1 negative imbalance. And DETR's set prediction, which removes NMS from the problem entirely rather than tuning it — the same move H6 makes one level up.

Reading: YOLO; Faster R-CNN; Focal Loss; DETR.

H2 — Segmentation, and averaging over the wrong pixels

Question. Skip connections are supposed to restore the fine detail that pooling destroys. How much are they worth, and where?

The math

The U-Net's claim is architectural. An encoder that pools discards high frequencies; upsampling cannot invent them back. A skip connection concatenates the pre-pool features into the decoder, giving the signal a route around the bottleneck. It is B4's residual argument in a second setting — the same idea one level up from a block.

Dice against cross-entropy is a question about weighting. Cross-entropy averages over pixels, so a class owning 3% of the image contributes 3% of the gradient. Soft Dice normalizes per class:

$$\mathcal{L}_{\text{Dice}} = 1 - \frac{1}{C}\sum_{c} \frac{2\sum_i p_{ic} y_{ic} + \epsilon}{\sum_i p_{ic} + \sum_i y_{ic} + \epsilon}$$

Each class contributes one term regardless of its pixel count. The hard Dice score uses an argmax and has zero gradient; replacing the indicator with the softmax probability p_{ic} is what makes it trainable. Since the *metric* (IoU per class) is itself area-normalized, the two objectives genuinely disagree about what "good" means on a background that owns 86% of the pixels.

IoU from the confusion matrix is the diagonal over the row plus column minus the double-counted diagonal, and Dice and IoU are monotone in each other, $D = 2I/(1 + I)$ — which the tests use to check one against the other.

What I built. A two-level U-Net on the same scenes H1 detects in. The no-skip arm keeps its entire decoder and loses only the concatenated channels, so the comparison is not confounded by capacity: 66,292 parameters against 59,812. Plus a boundary band — pixels within 2px of a label change, computed without scipy by disagreeing with a shifted copy and dilating — covering 11.7% of pixels.

Results (2000 train / 500 test, 10 epochs; medians over seeds 0–2). Pixel shares: background 0.856, rectangle 0.063, ellipse 0.048, triangle 0.032.

arm	foreground mIoU	boundary mIoU
skips + cross-entropy	0.9495	0.9237
no skips + cross-entropy	0.9378	0.8921
skips + Dice	**0.9573**	**0.9325**
skips + CE + Dice	0.9572	0.9327

	whole image	boundary band
skip gain	+0.0129	**+0.0299**

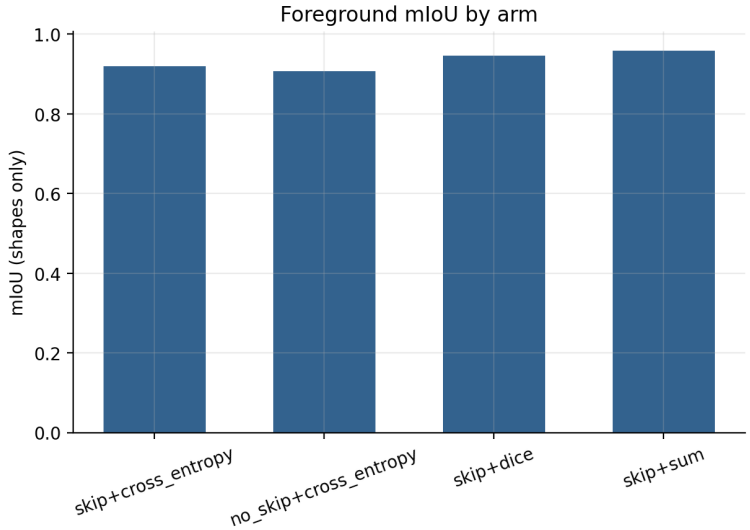


Figure 5: Arms

What surprised me. The first pass made the ablation look boring: +0.013 foreground mIoU for skips, everything above 0.94, a ceiling rather than a result. Then I scored the mechanism instead of the average. The argument for skip connections is about high frequencies — edges — and whole-image mIoU is dominated by object interiors, where a roughly-right blob already scores well. Restricting the *same predictions* to a 2px band around each true boundary gives **+0.030**

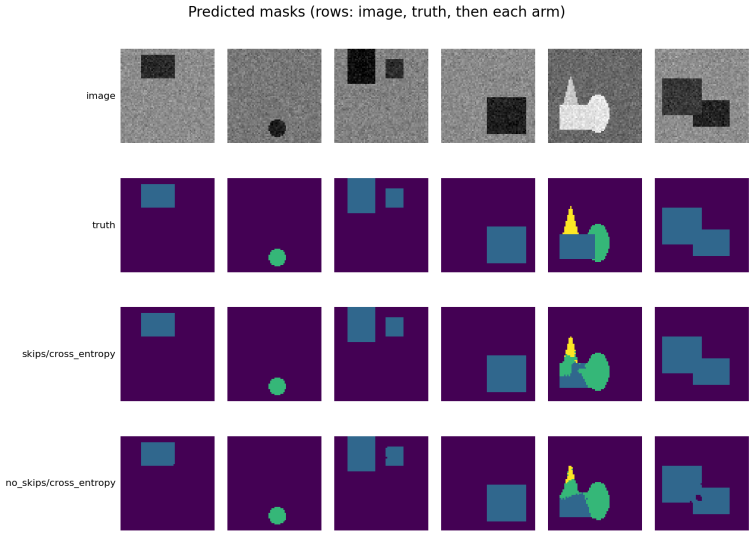


Figure 6: Masks

against +0.013, 2.3× larger, on all three seeds. The effect was always there; the metric was averaging the place skips act together with the place they don't.

And a claim that did not survive the sweep. On a single seed I recorded that Dice beats cross-entropy, which looked like the textbook class-imbalance story holding. Across three seeds `best/loss_is_dice` is $[0, 1]$ — Dice wins twice and loses once — over a loss spread of 0.012. Dice and CE+Dice are statistically indistinguishable here (0.9573 against 0.9572). I had reported the single-seed version as a finding before the sweep, and it was noise.

Open threads. A deeper U-Net, where more pooling should widen the skip gap. Thin structures rather than solid shapes, which is the regime the boundary result predicts skips matter most in. And boundary-weighted losses, which optimize directly for the band this rung had to invent to see the effect.

Reading: U-Net; Mask R-CNN; Segment Anything.

H3 — CTC, and what "held out" actually held out

Question. Reading text off an image has no alignment: a 128-pixel line holds an unknown number of characters at unknown positions. How much does removing that requirement cost, and where does the remaining error live?

The math

CTC marginalizes the alignment away. The network predicts a distribution over the alphabet plus a *blank* at every image column, and the loss sums the probability of every column labelling that collapses to the target:

$$p(y \mid x) = \sum_{\pi \in \mathcal{B}^{-1}(y)} \prod_{t} p(\pi_t \mid x),$$

where $\mathcal{B}$ deletes blanks and merges repeats. The blank is what makes repeated characters ex-

pressible at all — "AA" needs a blank between the two A's, or it collapses to "A" — and the sum is computed by a forward-backward recursion rather than enumerated.

Greedy decoding is not the most likely string. It takes the most likely *alignment* and collapses it, and those differ whenever one string's probability is spread over several alignments. Prefix beam search fixes that by scoring strings, which requires tracking two numbers per prefix — the probability it ends in a blank and the probability it ends in its final symbol — because which applies decides whether the next repeated character extends the prefix or merges into it. A test pins the case where they disagree: with two columns each splitting mass 0.55/0.45 between blank and 'A', greedy returns `""` and the beam returns `"A"`.

CER is an edit distance, not a per-position mismatch, because a model that drops one character would otherwise score as if everything after it were wrong too. It is pooled over the corpus rather than averaged per line, and it is not bounded by 1 — a hallucinating model can insert more than the reference holds.

What I built. A CRNN: convolutions pooling height to 1 while pooling width only twice, so a 32×128 image becomes 32 columns, then a two-layer BiLSTM and a per-column head. The width ratio carries a hard constraint — CTC cannot emit more characters than it has columns, and needs a spare column between repeats — so the rung checks it and exits loudly rather than discovering it as a silent `inf`.

Results (6000 train lines, 12 epochs; medians over seeds 0–2):

split	greedy CER	beam-15 CER	WER
seen fonts	0.0000	0.0000	0.0033
held-out fonts	0.1636	0.1621	0.4900

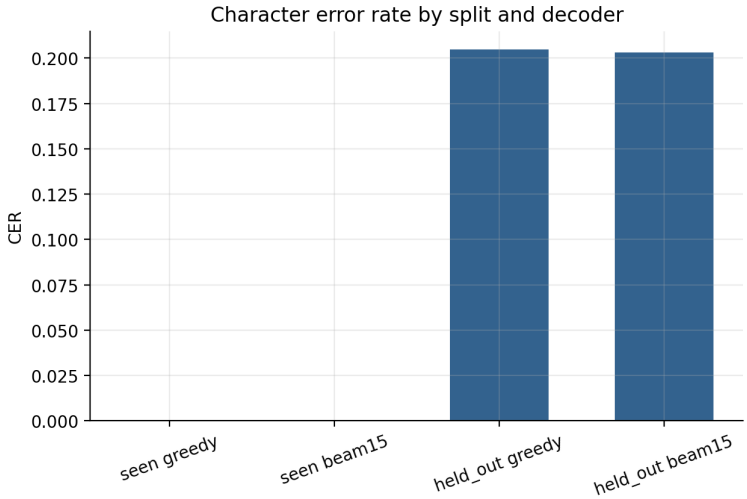


Figure 7: Decoders

What surprised me. Beam search buys **nothing**: 0.0000 CER at median against greedy, while held-out fonts cost 0.164. The textbook decoder upgrade, implemented properly with separate blank and symbol mass, is worth zero here, and the entire error budget is generalization. Decoding was never the bottleneck.

And the part-H finding I am least comfortable with. Breaking the held-out error down per font:

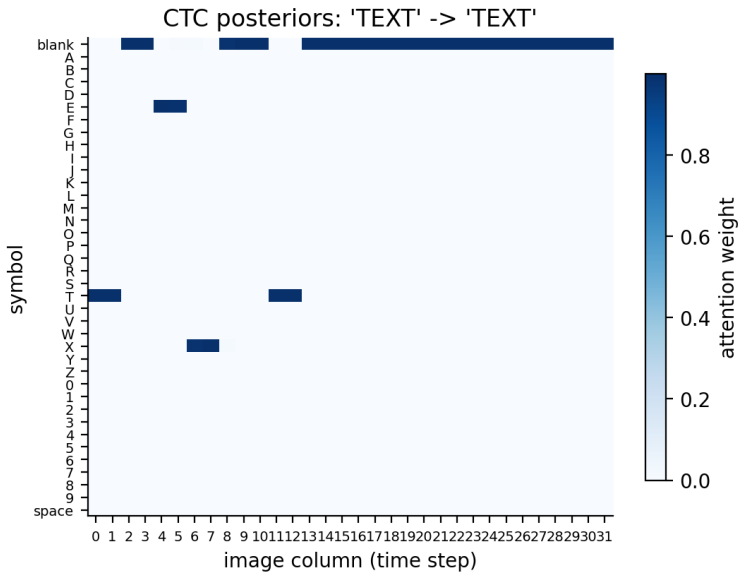


Figure 8: CTC posteriors

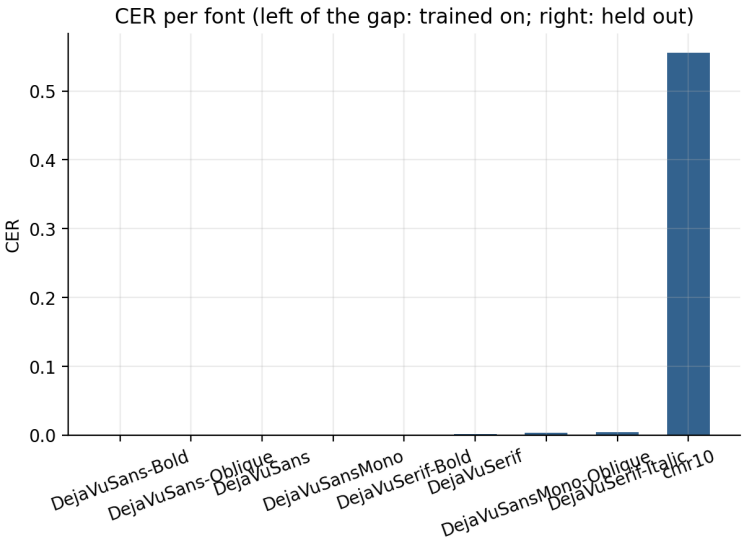


Figure 9: Error by font

held-out font	CER
DejaVuSansMono-Oblique	0.0038
DejaVuSerif-Italic	0.0052
cmr10 (Computer Modern)	**0.4666**

The "font gap" is one typeface. Two of my three held-out faces are read at 0.4–0.5% CER — essentially perfectly — because they are *DejaVu family members* whose siblings (DejaVuSansMono, DejaVuSerif) are in the training set. Only Computer Modern, a genuinely different design, is hard, and it carries the entire aggregate gap on its own.

So the honest statement is not "the model fails on unseen fonts" but "the model generalizes across weight and slant within a family and fails on an unfamiliar letterform" — which is a different and more useful claim. My held-out split was badly designed: I picked three faces for variety without checking whether they were near-duplicates of training faces, and the aggregate number hid that until the per-font breakdown went in. Every statement I made about this gap before the breakdown existed was an average over two easy cases and one hard one.

Open threads. A held-out set chosen by typeface family rather than by file, which is what would make the gap measure what I claimed. Font augmentation during training. And synthetic degradation — blur, skew, JPEG — which is where real OCR error comes from and which this data has none of.

Reading: CTC; CRNN; What Is Wrong With Scene Text Recognition Model Comparisons?.

H4 — Two ways to draw the image-to-string interface

Question. CTC assumes the output is monotonic in the input and conditionally independent given the columns. An attention decoder assumes neither and gains a language model over its own output. What does each assumption buy and cost?

The math

The independence assumption is the real difference, and it is visible in the factorization each model assumes:

$$p_{\text{CTC}}(\pi \mid x) = \prod_t p(\pi_t \mid x) \qquad \text{against} \qquad p_{\text{AR}}(y \mid x) = \prod_t p(y_t \mid y_{<t}, x).$$

CTC's columns are conditionally independent given the image, so it cannot condition on what it already wrote — it cannot spell. The autoregressive product carries $y_{<t}$, which is an implicit language model over the label distribution, obtained for free and never asked for.

The monotonicity assumption is the other. CTC's collapse is order-preserving by construction, so alignment is never something the model can get wrong. An attention decoder chooses where to look at every step and must *learn* both the alignment and when to stop — and an emitted EOS is a decision, not a structural consequence.

Attention is C2's, pointed at image columns rather than source tokens: $\alpha = \text{softmax}(h_t^\top H/\sqrt{d})$ with context $\sum_i \alpha_i H_i$, plus input feeding so the decoder knows where it just looked.

What I built. Three arms on byte-identical data. CTC is H3's CRNN, imported rather than reimplemented. The attention arm uses the *same* convolutional encoder and the *same* dot-product attention, so the comparison isolates the decoder. TrOCR is zero-shot, never trained here.

The experiment that matters is length: both arms train on one and two words and are tested on three. Critically there is a **two-words-on-the-wider-canvas control**, because at a fixed 128px a third word also shrinks the glyphs — and without separating those, a resolution failure reads as a length failure.

Results (medians over seeds 0–2; 0.4M / 0.9M / 62M parameters):

	CTC	attention	TrOCR
seen CER	0.0002	0.0000	0.0111
held-out CER	**0.1756**	0.1956	**0.0101**
held-out WER	0.4094	**0.2596**	0.0259
held-out exact	0.6125	**0.7013**	—
font gap	0.175	0.196	**−0.001**
length gap	**0.0003**	0.2202	—
width gap (control)	**−0.0002**	0.1393	—

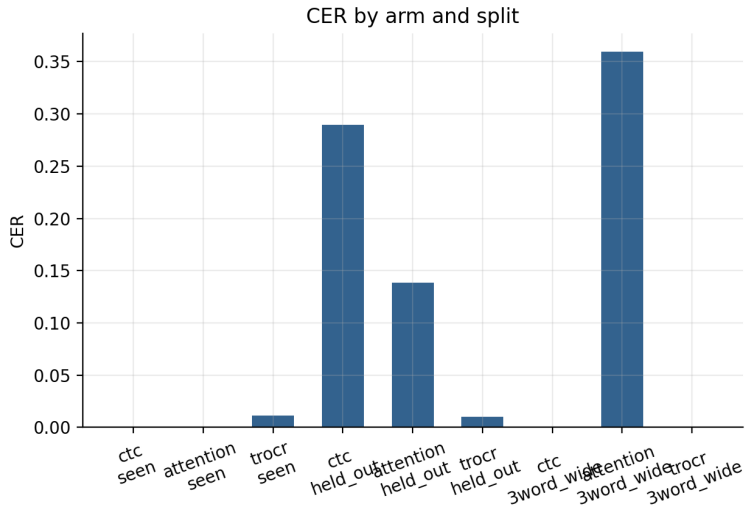


Figure 10: Arms

What surprised me. Three things.

CER and WER disagree about which decoder is better, on the same run. CTC wins on CER (0.176 vs 0.196); attention wins decisively on WER (0.260 vs 0.409) and exact match (0.701 vs 0.613). The CER→WER ratio is 2.5 for CTC and 1.3 for attention: the autoregressive decoder's implicit language model **concentrates** errors into fewer words, while CTC scatters independent character errors across many. Report CER and you ship CTC; report exact match and you ship attention. Nothing about the models differs between those two sentences.

The control did most of the work. The attention decoder's length gap is +0.220 — but **+0.139 of it is the wider canvas alone**, before a third word is involved. Only +0.081 is the unseen length. Had I run this without the control I would have written "attention fails to generalize to unseen length, +0.22 CER" and been 63% wrong about the cause. CTC meanwhile is *exactly* invariant to both: length gap +0.0003, width gap −0.0002. Its length bias at three words is −5.44 characters, i.e. it truncates, and it never gets a three-word line fully right.

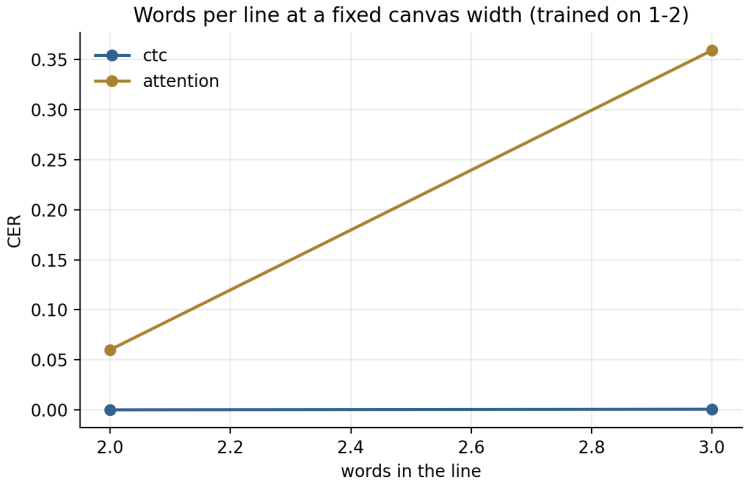


Figure 11: Length generalization

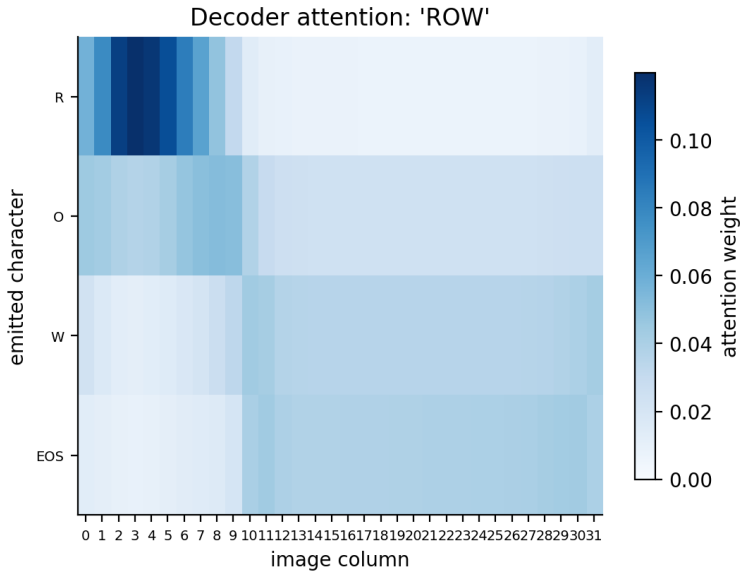


Figure 12: Attention map

TrOCR's font gap is negative. −0.001: it reads the held-out faces marginally *better* than the trained-on ones. Combined with H3's per-font breakdown, this settles the question — Computer Modern is not intrinsically hard, it was unfamiliar to a model trained on six faces. A model trained on real printed text has never heard of the problem.

Open threads. Beam search or a language model on the attention arm, which is where its implicit LM could be made explicit. Scheduled sampling for the exposure bias behind the truncation. And the obvious missing arm: fine-tuning TrOCR on this data rather than only reading it zero-shot.

Reading: TrOCR; Attention Is All You Need; Baek et al. on comparing recognizers honestly.

H5 — The pipeline, and where its error actually lives

Question. Detect then recognize is how every classical OCR system works, and it introduces a failure neither stage can exhibit alone: the second consumes the first's mistakes. Whose fault is the error?

The math

A single end-to-end number cannot answer that, so the same recognizer is measured three ways: on ground-truth crops (detection assumed perfect), on detected crops (the pipeline as deployed), and on ground-truth crops with each box edge jittered by a known amount (detection error simulated, detector removed from the loop). The third turns the gap between the first two from an observation into an explanation.

The jitter is per edge, not a shift. A uniform translation moves the crop but keeps the text framed inside it; independent edges change the *framing* — clipping a first letter, admitting a neighbour's descender — which is the failure a real detector produces.

The end-to-end metric is page-level CER over lines joined in reading order:

$$\mathrm{CER}_{\mathrm{page}} = \frac{\sum_p d\Big(\bigoplus_{i\in\sigma(p)} \hat{y}_i,\ \bigoplus_{i\in\sigma^\star(p)} y_i\Big)}{\sum_p \Big|\bigoplus_{i\in\sigma^\star(p)} y_i\Big|},$$

with $\oplus$ concatenation and σ the reading order induced by the predicted boxes. Writing it this way makes the choice explicit: because the prediction is a concatenation over *detected* lines and the reference over *true* ones, a missed line becomes a run of deletions, a spurious one a run of insertions, and a wrong σ costs the edit distance between two orderings. A per-line average would score only the lines that were found — exactly the population a broken detector biases.

What I built. H1's detector — its grid arithmetic generalized from square scenes to 160×256 pages, so H5 inherits H1's round-trip test rather than growing a private copy — feeding H3's CRNN, trained on ground-truth crops. Training the recognizer on clean crops and testing it on detected ones is the point: that distribution shift *is* the detector's error. Plus a second recognizer trained with ±4px jittered copies, because a diagnosis implies a remedy and the remedy should be measured.

Results (2500 train / 500 test pages, medians over seeds 0–2):

detector: recall 0.854 [0.611, 0.929] · mAP@0.5 0.730 [0.359, 0.807] · matched IoU 0.671 [0.523, 0.731] · 1.18 spurious boxes per page

	clean recognizer	+ crop augmentation
arm 1 — ground-truth crops	0.0001	0.0000
arm 2 — full pipeline	**0.5121**	**0.3670**
matched lines only	0.3640	0.1868
tolerance (5% CER)	±0.33 px	**±7.78 px**

detection's share of page error	**0.9998** [0.9996, 1.0000]

Jitter curves (line CER at 0/2/4/8/16 px) — clean: 0.0001, 0.305, 0.505, 0.709, 0.863. Augmented: 0.0000, 0.0001, 0.0003, 0.053, 0.427.

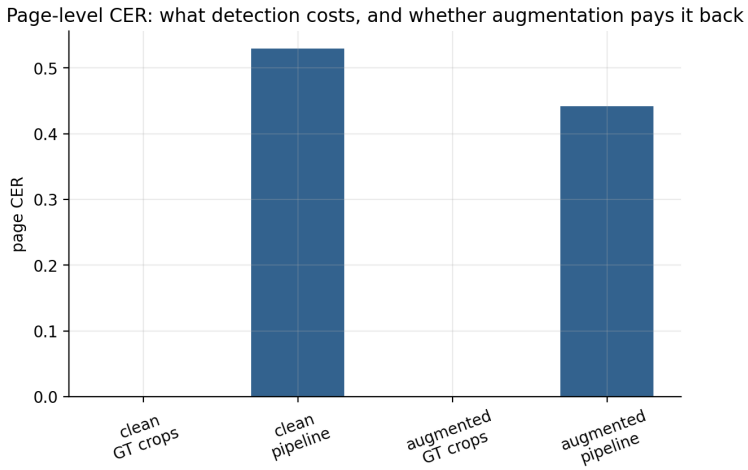


Figure 13: Attribution

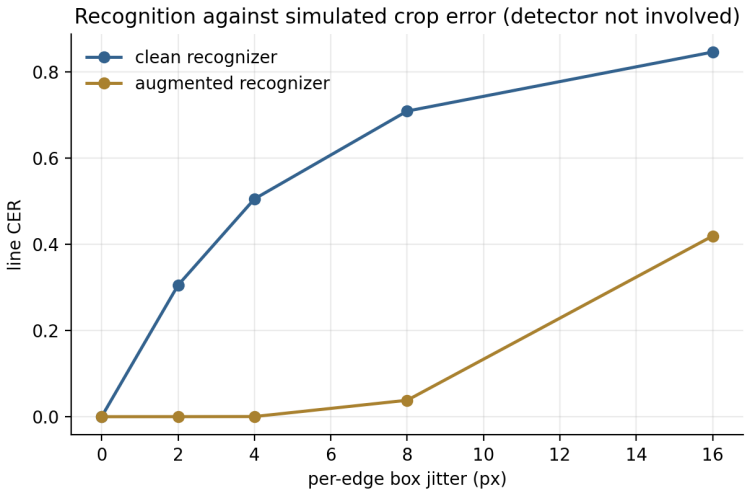


Figure 14: Jitter

What surprised me. Detection is 99.98% of the error on every seed. The recognizer on ground-truth crops is essentially perfect (0.0001 page CER) and the entire 0.512 appears the moment it is handed real boxes. A pipeline's quality is its weakest interface, not its weakest model.

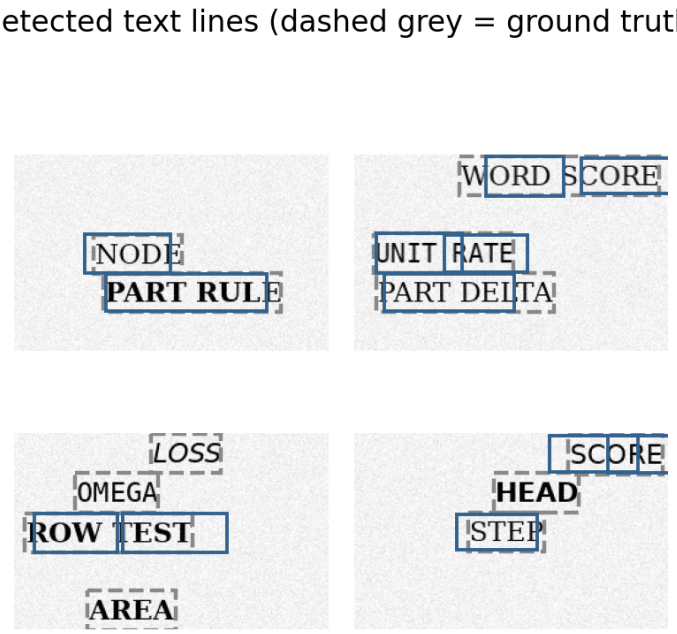


Figure 15: Pages

Arm 3 says why, and the number is brutal: interpolated, the clean recognizer crosses 5% CER at **0.33 pixels** of per-edge crop error. It was trained only on exact crops, has never seen a misframed line, and is brittle to a degree that makes the detector's 0.671 matched IoU fatal rather than merely mediocre.

The remedy worked, and mostly did not help. Training on jittered crops made the recognizer **24× more tolerant** ($\pm 0.33 \rightarrow \pm 7.78$ px) at *negative* cost on clean crops, and halved error on lines the detector found ($0.364 \rightarrow 0.187$). And the pipeline still sits at 0.367. Augmentation fixes framing; it cannot read a line that was never cropped. With recall at 0.854, the deletions from missed lines dominate everything crop quality controls. I had framed brittleness as *the* explanation for the pipeline gap — it is real, dramatic, and worth about a sixth of the total.

A single-seed number I reported and had to withdraw. The first H5 run gave recall 0.611 and mAP 0.359, and I described the detector's weakness using those. Across three seeds they are $[0.611, 0.929]$ and $[0.359, 0.807]$: I had quoted the *worst* seed as typical. The attribution conclusion is unaffected — 99.98% on all three — but the magnitude of the detector's failure is a range, not a point.

Open threads. Anchors with aspect priors: text lines are 54–189px wide and 32px tall, extreme ratios sharing one log-size regression, on a 20×32 grid that is 99.6% negative cells. Recall is the term that matters, so a lower score threshold and a recall-oriented detector would move the pipeline more than any recognizer change. And joint training through the crop, which is the last thing you can do before giving up on stages entirely.

Reading: EAST; CRAFT; PP-OCR on what a practical pipeline contains.

H6 — Deleting the interface

Question. H5 measured that the stage boundary holds 99.98% of the error. A vision-language OCR model removes the boundary rather than improving either stage. What does that cost?

The math

There is no box, so there is no box error. A page becomes visual tokens and the transcript is decoded autoregressively from them; detection, NMS, cropping and reading order all cease to exist as separately-failing components. The architectural question is what replaces them.

Token count is the cost driver, because attention is quadratic:

$$\text{FLOPs}_{\text{attn}} \approx 4LHT^2 d_h, \qquad T = \frac{HW}{P^2 m^2}$$

for patch size P and a spatial merge factor m. PaddleOCR-VL's 2×2 merge concatenates neighbouring patches before the projector, dividing T by four and therefore attention by **sixteen** — the cheapest available lever on a quadratic.

The KV cache is $2LH_{kv}d_h$ bytes per token per dtype byte, which is where grouped-query attention pays: 16 query heads against 2 key/value heads makes the cache 8× smaller at unchanged query capacity.

Mixture-of-experts decouples parameters held from parameters run. With E routed experts of width d_{moe} and top-k routing plus s shared experts, a layer holds $(E + s)$ experts' worth and runs $(k + s)$.

What I built. An accounting rung, not a benchmark. PaddleOCR-VL's architecture is **parsed out of candle's `config.rs` at run time** rather than quoted, so a fork that moves is detected instead of silently misdescribed; the 17 parsed defaults independently match the published `config.json`. The parameter formulas are validated against a model the rung can actually load and count — TrOCR, predicted 60.8M against measured 61.6M, **1.3% off** — so the budgets below are a checked calculation rather than an assertion.

Results:

component	parameters
PaddleOCR-VL vision (NaViT, 27 layers, 1152 wide)	411.4M
PaddleOCR-VL projector (2×2 merge + MLP)	26.0M
PaddleOCR-VL text (ERNIE-4.5, 18 layers, GQA 16:2)	466.7M
PaddleOCR-VL total	**904.0M** (45.5% vision)
Unlimited-OCR SAM-ViT-B tower	85.0M
Unlimited-OCR CLIP-L tower	302.3M
Unlimited-OCR text (MoE, 64+2 experts, top-6)	**2.93B held / 740M active**

page	patches	tokens after 2×2 merge	attention
384px	729	182	4.1 GFLOP
1024px	5,329	1,332	220.7 GFLOP
A4 at 300dpi	44,250	11,062	—

KV cache 18.0 KiB/token, 8× less than multi-head would need.

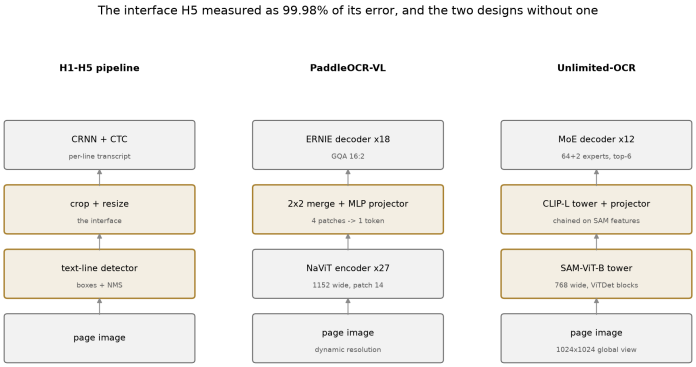


Figure 16: Architectures

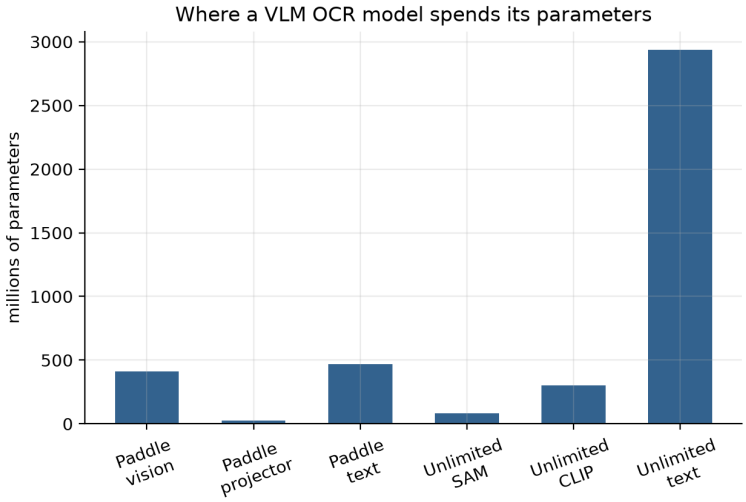


Figure 17: Parameter budget

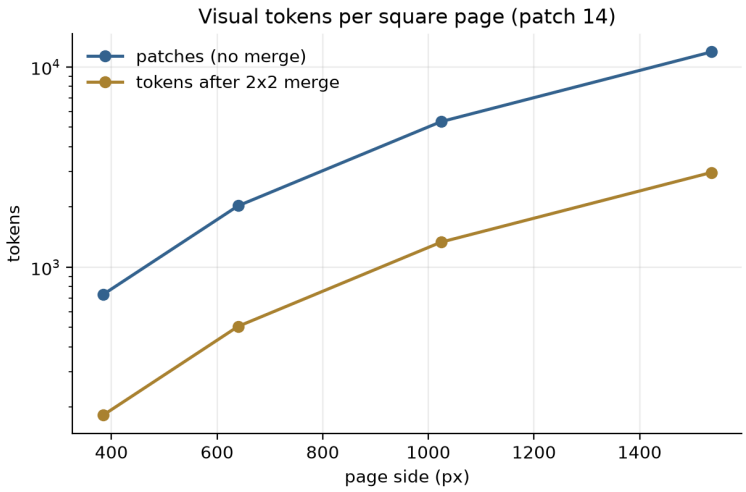


Figure 18: Token cost

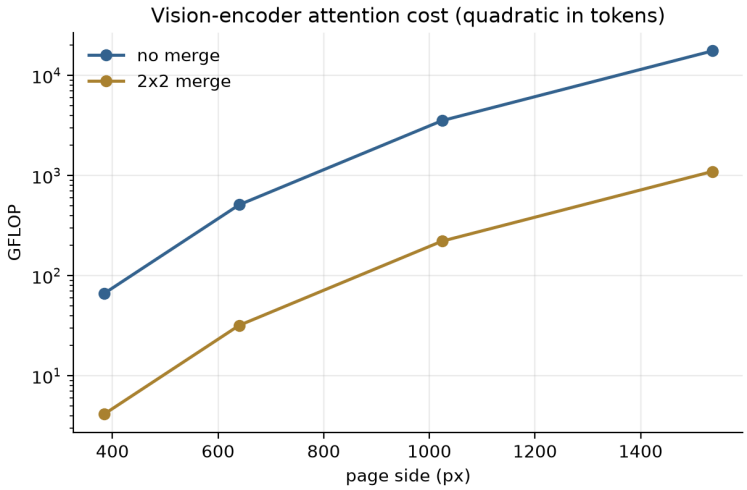


Figure 19: Attention cost

What surprised me. The vision tower is **45.5%** of PaddleOCR-VL — nearly as large as the language model it feeds. I had assumed the encoder would be a thin front-end on a big decoder, the way TrOCR's DeiT-small is. For document understanding the reading *is* the hard part, and the parameter split says so.

The exchange rate against H5 is **1,853×** — 904M parameters against 488k — for a model that also handles 109 languages, tables and formulas, so the comparison is not clean. But the structural point survives it: every component H5 measured as its error source is absent from that 904M, not improved within it.

And an error I made worth recording. I first counted Unlimited-OCR's text tower as a dense stack, which understated it by 4× and erased the distinction MoE exists to create. I had also *guessed* its `intermediate_size` at 6848 because candle's config carries no defaults for it — the guess happened to be right, which is worse than being wrong, because it would have shipped ungrounded and looked correct. Both are now read from the published config with provenance recorded beside them.

Open threads. Actually running PaddleOCR-VL on H5's pages, which would turn this accounting into the head-to-head the rung deliberately does not attempt. Measuring the vision tower's activation memory, which at 11,062 tokens for an A4 scan is likely the real deployment constraint rather than parameters. And DeepSeek-OCR's compression claims, which is the lineage Unlimited-OCR comes from.

Reading: PaddleOCR-VL; NaViT; Donut; GQA.

What part H settled

1. **A pipeline's quality is its weakest interface, not its weakest model.** H5's recognizer is perfect on clean crops and the pipeline is at 0.512 page CER, with 99.98% of that attributable to detection on every seed. Improving the recognizer 24× in robustness — which crop augmentation genuinely did, for free — recovered a sixth of it. The rest is lines that were never found.
2. **Detection has no metric until you choose one, and the choice is worth more than most modelling decisions.** The same predictions score 0.955 or 0.650 depending only

on what overlap counts as correct, and the NMS dial that moves mAP@0.5 by 0.109 moves mAP@0.75 by 0.032. Reporting a single mAP without its threshold conveys almost nothing.

3. **Which decoder is "better" depends on which metric you report.** On identical held-out data CTC wins CER and the attention decoder wins WER and exact match, because an autoregressive decoder's implicit language model concentrates errors into fewer words while CTC scatters them. Both statements are true of the same run.
4. **Averaging hides mechanisms; score where the mechanism acts.** Skip connections looked worth +0.013 until the same predictions were scored on a 2px boundary band, where they are worth +0.030. The effect was never absent — whole-image mIoU was averaging the place skips work together with the place they don't.
5. **Controls and breakdowns changed three conclusions that a headline number got wrong.** H4's canvas-width control showed 63% of an apparent length-generalization failure was resolution. H3's per-font breakdown showed the entire "held-out font gap" was one typeface, with the other two held-out faces read at 0.4% CER because they were family members of training faces. H1's localization investigation ruled out three plausible explanations and left me with an honest "I don't know" for the per-class ordering. In each case the aggregate was not wrong, it was uninformative in a way that looked informative.
6. **The single-seed habit keeps costing me.** H2's "Dice beats cross-entropy" and H5's "recall 0.611" were both reported before the sweep and both failed it — the first was noise over a 0.012 spread, the second was the worst of three seeds quoted as typical. Together with part B's `usable_lrs`, part D's `augmentation_gain` and H5's own degenerate tolerance metric, that is six occasions in this repo where a number that looked like a result was an artifact of how it was computed or how few times it was run. The arm-level measurements were right every time; only the summaries lied.
7. **The field's answer to H5 was to delete the interface, not fix it.** H6 prices that at 904M parameters and 182 tokens per 384px page. Whether that is worth it is a deployment question, but the direction is not mysterious once you have measured that the interface was where all the error was.

Cost of the part: ~3,400 lines including tests, 76 test cases, and a three-seed sweep in about 24 minutes on a T4 (H1 14s, H2 42s, H3 116s, H4 114s, H5 194s, H6 1s per seed). H5 dominates because it trains a detector and two recognizers; H6 is arithmetic.

Figures regenerate with `uv run python tools/writeup_figures.py H`; the numbers they are quoted against live in `data/h-metrics.json`. The part README is the short version of this document.